

Artificial Intelligence

Lecture 2c: Adversarial Search

Luís Paulo Reis

lpreis@fe.up.pt

Director of LIACC – Artificial Intelligence and Computer Science Lab.
Associate Professor at DEI/FEUP – Informatics Engineering Department,
Faculty of Engineering of the University of Porto, Portugal
President of APPIA – Portuguese Association for Artificial Intelligence



Games as Search Problems

- Hostile Agent (opponent) included in the world!
- Unpredictable Opponent => Solution is a Contingency Plan
- Time Limit => Unlikely to find goal! Approximation is necessary
- One of the oldest areas of AI! In 1950 Shannon and Turing created the first Chess programs!
- Chess:
 - Everyone believes that intelligence is needed to play
 - Simple rules but the game is complex
 - World fully accessible to the agent
 - Average branching factor of 35, 50-game play => 35^{100} leaves in the search tree (although there are only 10^{40} legal positions)
- Cutting concepts in the search tree and evaluation function!

Game Types

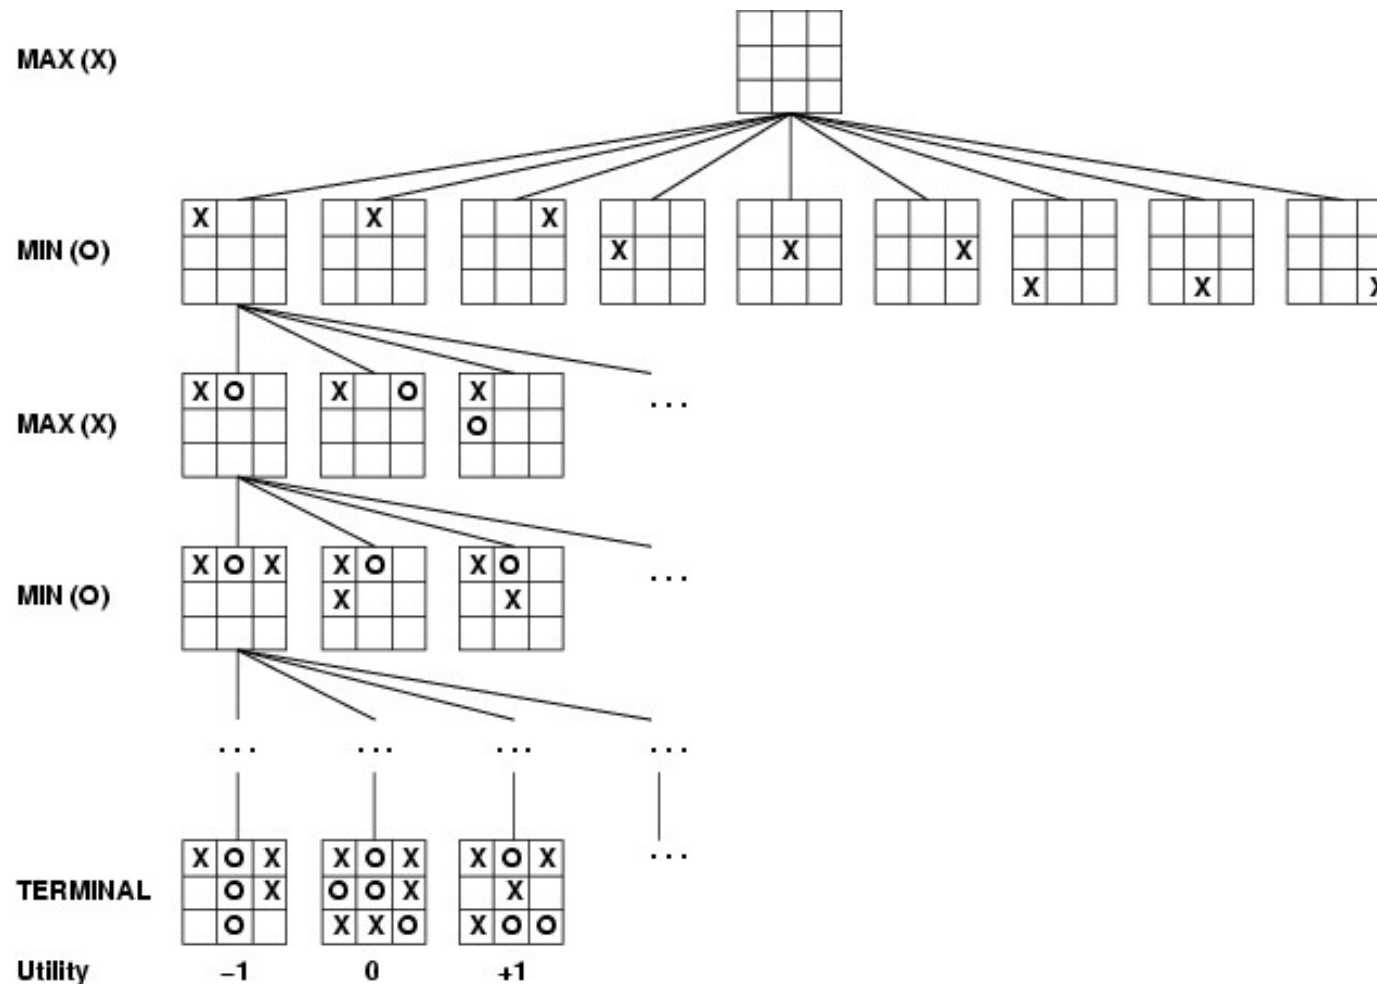
- **Information:**
 - Perfect: Chess, Checkers, Go, Othello, Backgammon, Monopoly
 - Imperfect: Poker, Scrabble, Bridge, King
- **Luck / Deterministic:**
 - Deterministic: Chess, Checkers, Go, Othello
 - Games with luck: Backgammon, Monopoly, Poker, Scrabble, Bridge, King
- **“Attack” plan:**
 - Algorithm for the perfect game
 - Finite horizon, approximate evaluation
 - Tree cuts to reduce costs

Perfect Decisions in Opponent Games - MiniMax Algorithm

- **Game: Search problem with:**
 - Initial State (board position and who is next to play)
 - Set of Operators (which define the legal movements)
 - Terminal Test (which determines whether the game is over or is in a terminal state)
 - Utility function (which gives a numerical value to the result of the game, for example 1-win, 0-draw, -1-defeat)
- **Minimax algorithm strategy:**
 - Generate the complete tree to the end states
 - Apply the utility function to these states
 - Calculate the utility values back to the root of the tree, one layer at a time
 - Choose the movement with the highest value!


Minimax – Tic-Tac-Toe Game Example

Easy to solve the Tic-Tac-Toe Game using Minimax




Minimax – General Example

- **Strategy:**
 - Choose the movement that has the highest minimax value = the best that can be achieved against the best responses from the opponent!



Alpha-Beta Cuts - Example

The grayed-out subtrees don't need to be explored (when moves are evaluated from left to right), since it is known that the group of subtrees as a whole yields the value of an equivalent subtree or worse, and as such cannot influence the final result. The Max and Min levels represent the turn of the player and the adversary, respectively



Minimax Algorithm

function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

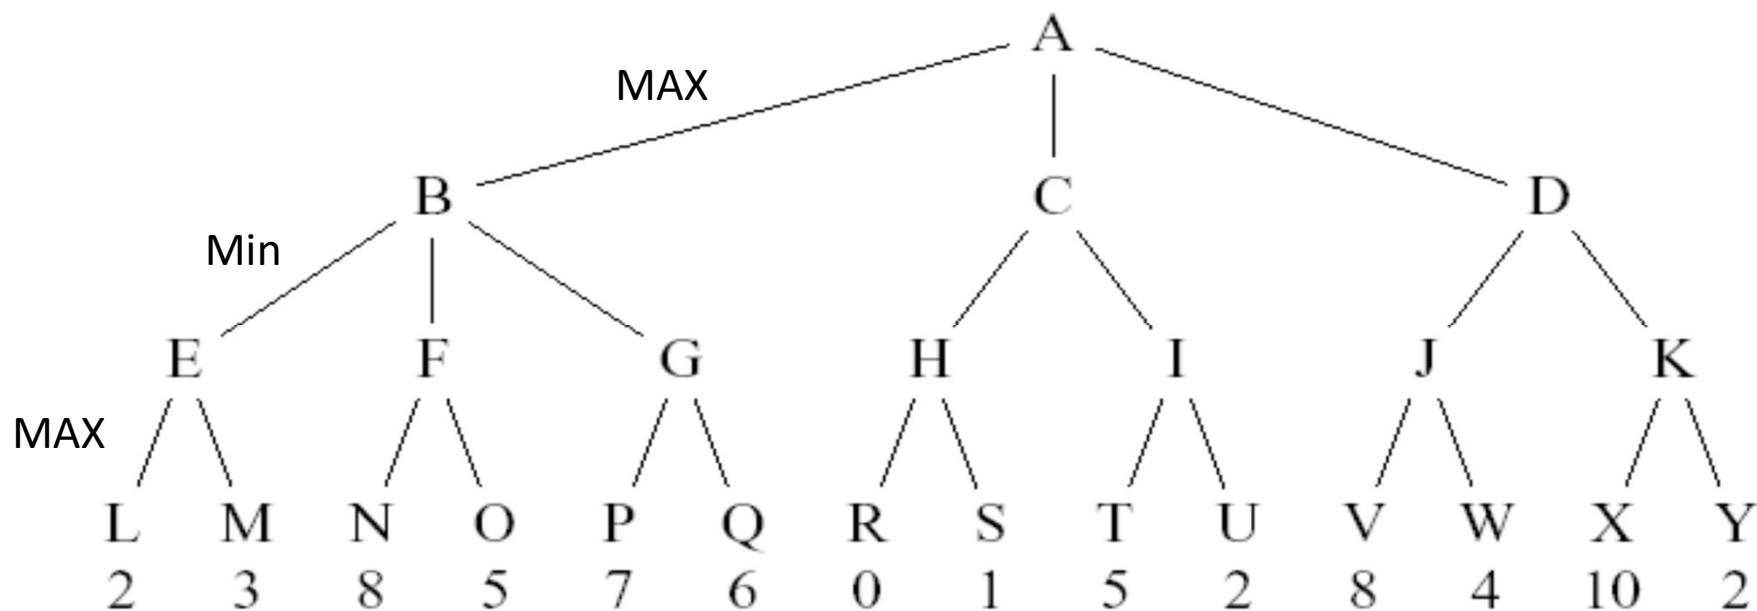
for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*


Exercise – Minimax Algorithm

- Assuming that MAX is the first to play, apply the Minimax Algorithm to the following tree, indicating the movement selected by the algorithm and the respective estimated value



Exercise – Minimax Algorithm

- Assuming that MAX is the first to play, apply the Minimax Algorithm to the following tree, indicating the movement selected by the algorithm and the respective estimated value.
- Solution: A → D (expected value = 8, e.g. oponente will play D→J and we will play J→V)




Minimax Properties

- **Properties:**
 - Complete? Yes if the tree is finite!
 - Optimal? Yes against an optimal opponent! If no?
 - Time Complexity? $O(b^m)$
 - Complexity in Space? $O(bm)$ (exploration first in depth)
- **Problem:**
 - Not viable for any minimally complex game
- **Example:**
 - For chess ($b = 35$, $m = 100$), $b^m = 35^{100} = 2.5 * 10^{154}$
 - Assuming that 450 million hypotheses are analysed per second $\Rightarrow 2 * 10^{138}$ years to arrive at the solution!


Limited Resources

- Assuming we have 100 seconds and explore 10^4 nodes / second, we can explore 10^6 nodes per movement
- Usual approach:
 - Cutting Test: Depth Limit
 - Evaluation Function: Estimated utility (interest) for the position/state
- Example (Chess):
 - Cutting Test: Depth of Analysis n
 - Simple evaluation function = sum of the values of the white pieces in play minus the sum of the values of the black pieces in play!
 - Evaluation function should only be applied to stable positions (in terms of their value). For example, positions with possible captures should be further explored ...
- Another Problem: Horizon problem!

Chess – Evaluation Functions




(a) White to move



(b) White to move

Chess – Horizon Problem

- With a limited horizon it seems interesting for Black to check with the pawns and delay the capture of the bishop (beyond the horizon)









Cuts - *MinimaxCutoff*

- ***MinimaxCutoff* is identical to *MinimaxValue* except:**
 - Terminal-Test is replaced by *Cutoff*
 - Utility is replaced by Evaluation (which calculates an evaluation of the reached position)
- **Does it work in practice?**
 - If $b^m = 10^6$ with $b = 35 \Rightarrow m = 4$
- **Playing chess with depth 4 is absolutely miserable!**
 - Depth 4 \Rightarrow Newbie Player
 - Depth 8 \Rightarrow PC, Good Human Player
 - Depth 12 \Rightarrow Deep Blue, Kasparov

Alpha-Beta Cuts (1)

- α is the best value (for Max) found so far on the current path
- If V is worse than α , Max should avoid it \Rightarrow cut off the branch
- β is defined in the same way for Min




Alpha-Beta Cuts (2)

- **Alpha-Beta cuts do not affect the final result**
- **Good ordering improves cutting efficiency**
 - Essential to reason about node ordering
- **Perfectly arranged: Time Complexity = $O(b^{m/2})$**
 - Doubles search depth
 - Depth 8 => Good chess player
- **Good example of the value of reasoning about which computations are relevant:**
 - This is essential in Intelligent Systems


Exercise – MINIMAX with Cuts

- Assuming that MAX is the first to play, apply the Minimax Algorithm with Alpha-Beta cuts to the following tree, indicating the movement selected by the algorithm. Indicate graphically and justify all the cuts made when applying the Minimax algorithm



Exercise – MINIMAX with Cuts

- Assuming that MAX is the first to play, apply the Minimax Algorithm with Alpha-Beta cuts to the following tree, indicating the movement selected by the algorithm. Indicate graphically and justify all the cuts made when applying the Minimax algorithm
- Solution:



Exercise – MINIMAX with Cuts

- Assuming that MAX is the first to play, apply the Minimax Algorithm with Alpha-Beta cuts to a tree with three levels, a branch factor 3 and with the following values of the evaluation function for the final line:

[8 6 3 1 10 3 15 20 6 7 4 6 25 10 4 4 3 5 1 20 4 12 1 10 4 22 10]

- Indicate graphically and justify all the cuts made when applying the Minimax algorithm

Negamax

- Negamax search is a variant form of minimax search that relies on the zero-sum property of a two-player game.
- This algorithm relies on the fact that
$$\max(a,b) = -\min(-a,-b)$$
- to simplify the implementation of the minimax algorithm.
- More precisely, the value of a position to player A in such a game is the negation of the value to player B. Thus, the player on move looks for a move that maximizes the negation of the value resulting from the move: this successor position must by definition have been valued by the opponent.
- The reasoning of the previous sentence works regardless of whether A or B is on move. This means that a single procedure can be used to value both positions.
- This is a coding simplification over minimax, which requires that A selects the move with the maximum-valued successor while B selects the move with the minimum-valued successor.

Negamax - Pseudocode

```
function negamax(node, depth, color) is  
    if depth = 0 or node is a terminal node then  
        return color × the heuristic value of node  
    value :=  $-\infty$   
    for each child of node do  
        value := max(value, -negamax(child, depth - 1, -color))  
    return value
```

```
(* Initial call for Player A's root node *)  
negamax(rootNode, depth, 1)
```

```
(* Initial call for Player B's root node *)  
negamax(rootNode, depth, -1)
```

Negamax Alpha-Beta - Pseudocode

```
function negamax(node, depth,  $\alpha$ ,  $\beta$ , color) is  
    if depth = 0 or node is a terminal node then  
        return color  $\times$  the heuristic value of node  
  
    childNodes := generateMoves(node)  
    childNodes := orderMoves(childNodes)  
    value :=  $-\infty$   
    foreach child in childNodes do  
        value := max(value, -negamax(child, depth - 1,  $-\beta$ ,  $-\alpha$ , -color))  
         $\alpha$  := max( $\alpha$ , value)  
        if  $\alpha \geq \beta$  then  
            break (* cut-off *)  
    return value
```

```
(* Initial call for Player A's root node *)  
negamax(rootNode, depth,  $-\infty$ ,  $+\infty$ , 1)
```

Deterministic Games

- **Checkers:**
 - **Chinook ended the 40-year reign of human champion Marion Tinsley in 1994.** It used a database for match ends defining the perfect way to win for all positions involving 8 or less pieces (in total 443748401247 positions). Nowadays checkers is a solved problem.
- **Chess:**
 - **Deep Blue defeated human world champion Gary Kasparov** in a 6-matches game **in 1997**. Deep Blue searched 200 million positions per second and used an extremely sophisticated evaluation function and (undisclosed) methods to extend some lines of search beyond depth 40!
- **Othello:**
 - **Human champions refuse to compete with computers** because they have no chance! (b between 5 and 15)

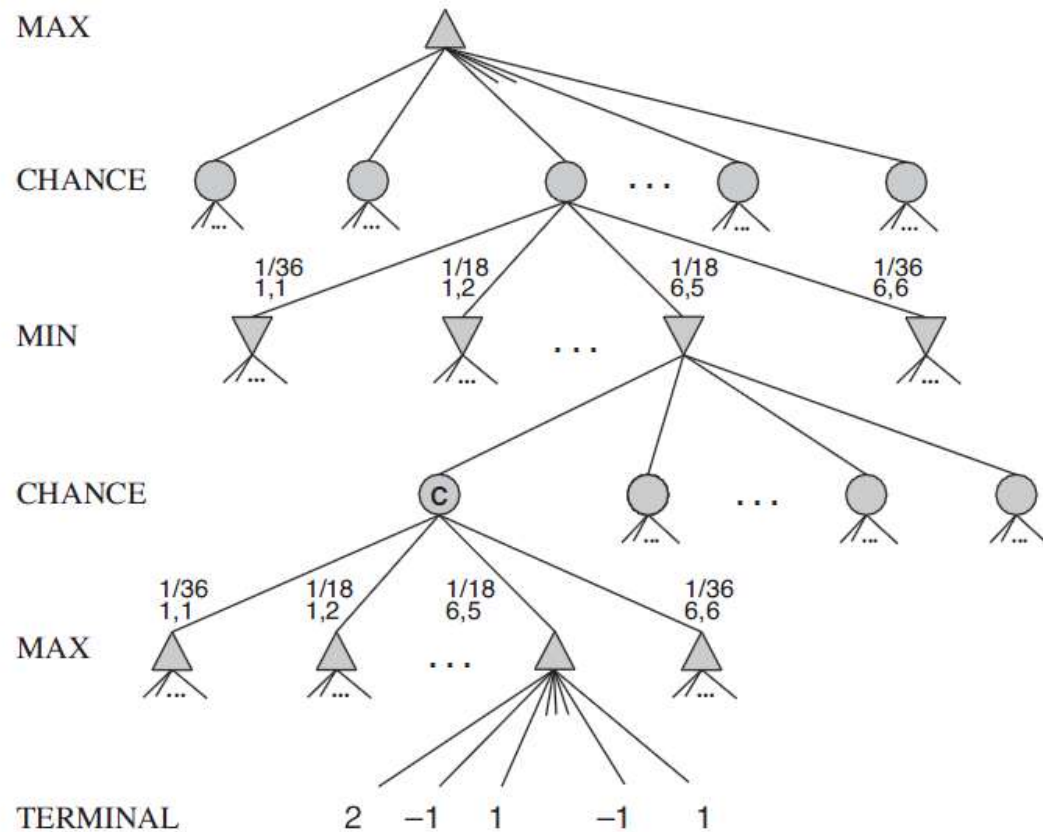
Deterministic Games - Go

- **Go (2015):**
 - Human champions refuse to compete with computers because machines cannot play reasonably ($b > 300$)
- **Go (2017-2020):**
 - AlphaGo (Fan, Lee), AlphaGo Master, AlphaGo Zero and AlphaZero!
Machines beat 100-0 human champions and previous machines.



Games of Luck - *ExpectiMiniMax*

- In many games, unlike chess, there are external events that affect the game, such as drawing a card or rolling a dice!
- Examples: Card games, Backgammon, Scrabble, ...
- Search tree must include probability nodes!
- Decision is made based on the expected value!
- *ExpectiMiniMax*



ExpectiMiniMax - Pseudocode

```
function expectiminimax(node, depth)
  if node is a terminal node or depth = 0
    return the heuristic value of node
  if the adversary is to play at node
    // Return value of minimum-valued child node
    let  $\alpha := +\infty$ 
    foreach child of node
       $\alpha := \min(\alpha, \text{expectiminimax}(\text{child}, \text{depth}-1))$ 
  else if we are to play at node
    // Return value of maximum-valued child node
    let  $\alpha := -\infty$ 
    foreach child of node
       $\alpha := \max(\alpha, \text{expectiminimax}(\text{child}, \text{depth}-1))$ 
  else if random event at node
    // Return weighted average of all child nodes' values
    let  $\alpha := 0$ 
    foreach child of node
       $\alpha := \alpha + (\text{Probability}[\text{child}] \times \text{expectiminimax}(\text{child}, \text{depth}-1))$ 
  return  $\alpha$ 
```

Exercises – Adversarial Search/Games

Formulate some of the following 2-player board games as a adversarial search problem (game) and design an intelligent agent capable of playing it, using Minimax with alpha-beta cuts:

- | | | |
|---|---|--|
| 1. Chess | 19. Link Five | 35. Paper & Pencil Racing |
| 2. Shogi | 20. Mancala (normal/4x8) | 36. Arimaa |
| 3. Checkers | 21. Fanorana | 37. Gipf |
| 4. Connect 4 | 22. Nine Mens Morris | 38. Lines of Action |
| 5. Attaxx | 23. Alquerque | 39. Connections |
| 6. Chinese Checkers | 24. Tablut | 40. Fanorama |
| 7. Othello (Reversi) | 25. Surakarta | 41. Hexxagon |
| 8. Abalone | 26. Terrace | 42. Jungle Game |
| 9. Hex | 27. Go | 43. Seega |
| 10. Omega Chess | 28. Dots and Boxes | 44. Halma (2 players) |
| 11. Tori Shogi | 29. Dots and Hexagons | 45. Quits (Gigamic) |
| 12. Quoridor (Normal/Kids)
(Gigamic) | 30. Amazons | 46. Pylos (Gigamic) |
| 13. Quarto (Gigamic) | 31. Scrabble (visible letters and
known draws) | 47. Tantrix (2 players) |
| 14. Quixo (Gigamic) | 32. Tic tac toe (normal/ memory/
movement, 3D) | 48. Carcassone (2 players, known
draws) |
| 15. Quads (Gigamic) | 33. Dominoes (visible pieces and
known draws) | 49. Blokus |
| 16. Sahara (Gigamic) | 34. Backgammon (known draws) | 50. Sputnik (Gigamic) |
| 17. Pentaminoes (8x8) | | 51. Katamino (Gigamic) |
| 18. Stratego | | 52. Gobblet (Gigamic) |

Summary - Games

- **Working with games is extremely interesting (but also dangerous...)**
 - Easy to test new ideas!
 - Easy to compare agents with other agents
 - Easy to compare agents with humans!
- **Games illustrate several interesting points of AI:**
 - Perfection is unattainable => it is necessary to approximate!
 - It is a good idea to think about what to think!
 - Uncertainty restricts the allocation of values to states!
- **Games work for AI like Formula 1 for car building!**

Artificial Intelligence

Lecture 2c: Adversarial Search

Luís Paulo Reis

lpreis@fe.up.pt

Director of LIACC – Artificial Intelligence and Computer Science Lab.
Associate Professor at DEI/FEUP – Informatics Engineering Department,
Faculty of Engineering of the University of Porto, Portugal
President of APPIA – Portuguese Association for Artificial Intelligence

